

Efficient Decentralized Monitoring of Safety in Distributed Systems

Koushik Sen, Abhay Vardhan, Gul Agha, Grigore Roşu

Department of Computer Science

University of Illinois at Urbana Champaign

{ksen, vardhan, agha, grosu}@cs.uiuc.edu

Abstract

We describe an efficient decentralized monitoring algorithm that monitors a distributed program's execution to check for violations of safety properties. The monitoring is based on formulae written in PT-DTL, a variant of past time linear temporal logic that we define. PT-DTL is suitable for expressing temporal properties of distributed systems. Specifically, the formulae of PT-DTL are relative to a particular process and are interpreted over a projection of the trace of global states that represents what that process is aware of. A formula relative to one process may refer to other processes' local states through remote expressions and remote formulae. In order to correctly evaluate remote expressions, we introduce the notion of KNOWLEDGEVECTOR and provide an algorithm which keeps a process aware of other processes' local states that can affect the validity of a monitored PT-DTL formula. Both the logic and the monitoring algorithm are illustrated through a number of examples. Finally, we describe our implementation of the algorithm in a tool called DIANA.

1. Introduction

Software errors from a number of different problems such as incorrect or incomplete specifications, coding errors, and faults and failures in the hardware, operating system or network. Model checking is an important technology which is finding increasing use as a means of reducing software errors. Unfortunately, despite impressive recent advances, the size of systems for which model checking is feasible remains rather limited. This weakness is particularly critical in the context of distributed systems: concurrency and asynchrony results in inherent non-determinism that significantly increases the number of states to be analyzed. As a result, most system builders must continue to use testing to identify bugs in their implementations.

There are two problems with software testing. First, testing is generally done in an *ad hoc* manner: the software developer must hand translate the requirements into specific dynamic checks on the program state. Second, test coverage

is often rather limited, covering only some execution paths. To mitigate the first problem, software often includes dynamic checks on a system's state in order to identify problems at run-time. Recently, there has been some interest in run-time monitoring techniques which provide a little more rigor in testing. In this approach, monitors are automatically synthesized from a formal specification. These monitors may then be deployed off-line for debugging or on-line for dynamically checking that safety properties are not being violated during system execution.

In this paper, we argue that distributed systems may be effectively monitored at runtime against formally specified safety requirements. By effective monitoring we mean not only linear efficiency, but also decentralized monitoring where few or no additional messages need to be passed for monitoring purposes. We introduce an epistemic temporal logic for distributed knowledge. We illustrate the expressiveness of this logic by means of some simple examples. We then show how efficient distributed monitors may be synthesized from the specified requirements. Finally, we describe a distributed systems application development framework, called DIANA. To use DIANA, a user must provide an application together with the formal safety properties that she wants monitored. DIANA automatically synthesizes code for monitoring the specified requirements and weaves appropriate instrumentation code into the given application. The architecture of DIANA is illustrated in Figure 1.

The work presented in this paper was stimulated by the observation that in many distributed systems, such as wireless sensor networks, it is quite impractical to monitor requirements expressed in classical temporal logics. For example, consider a system of mobile nodes in which one mobile node may request a certain value from another mobile node. On receiving the request, the second node computes the value and returns it. An important requirement in such a system is that no node receives a reply from a node to which it has not previously issued a request. It is easy to see that Linear Temporal Logic (LTL) would not be a practical specification language for any reasonably sized collec-

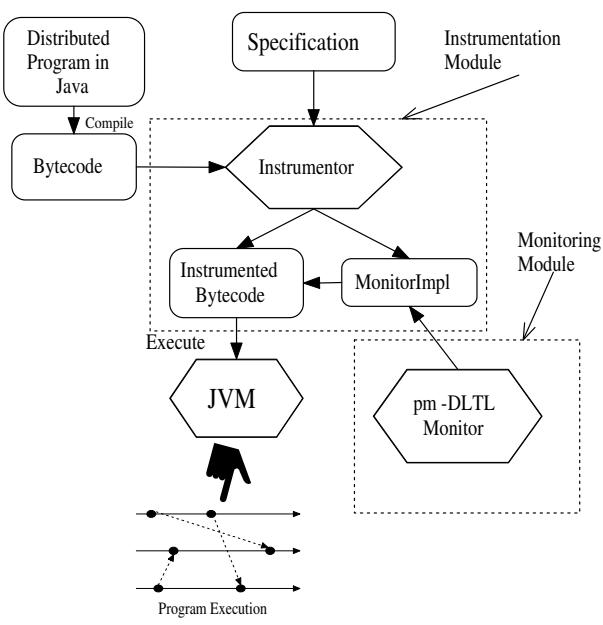


Figure 1. The Architecture of DiANA

tion of nodes. To use LTL, we would need to collect consistent snapshots of the global system; a monitor would then check the snapshots for possible violations of the property by considering all possible interleavings of events that are allowed by the distributed computation. In a system of thousands of nodes, collecting such a global snapshot would be prohibitive. Moreover, the number of possible interleavings to be considered would be large even if powerful techniques such as partial order reduction are used.

To address the above difficulty, we define *past-time distributed temporal logic* (PT-DTL). Using PT-DTL, one can check a property such as the one above by having a local monitor on each node. For example, node a monitors “if a has received a value then it must be the case that previously in the past at b the following held: b has computed the value and at a a request was made for that value in the past”. This is precisely and concisely expressed by the PT-DTL formula:

$$\text{receivedValue} \rightarrow \\ @_b(\Diamond(\text{computedValue} \wedge @_a(\Diamond\text{requestedValue})))$$

Note that we read $@$ as “at”, $@_b F$ is the value of F in the most recent local state of b that the current process is aware of, and \Diamond denotes the formula was true sometimes in the past. Monitoring the above formula involves sending no additional messages – it involves inserting only a few bits of information which are piggybacked on the messages that are already being passed in the computation. This efficiency provides a substantial improvement over what is required to monitor formulas written in classical LTL.

We introduce *remote expressions* in PT-DTL to represent values which are functions depending on the state of a remote process. For example, a process may monitor the property: “if my alarm has been set then it must be the case that the difference between my temperature and the temperature at process b exceeded the allowed value”. This is expressed as:

$$\text{alarm} \rightarrow \Diamond((\text{myTemp} - @_b\text{otherTemp}) > \text{allowed})$$

Here $@_b\text{otherTemp}$ is a remote expression that is subtracted from the local value of myTemp .

An example of a safety property that may be useful in the context of an airplane software is: “if my airplane is landing then the runway allocated by the airport matches the one that I am planning to use”. This property may be expressed in PT-DTL as follows:

$$\text{landing} \rightarrow (\text{runway} = (@_{\text{airport}}\text{allocRunway}))$$

Many researchers have proposed temporal logics to reason about distributed systems. Most of these logics are inspired by the classic work of Aumann [5] and Halpern *et al.* [7] on knowledge in distributed systems. Meenakshi *et al.* define a knowledge temporal logic interpreted over a message sequence charts in a distributed system [16] and develop methods for model checking formulae in this logic. Our communication primitive was in part inspired by this work, but we allow arbitrary expressions and atomic propositions over expressions in their logic.

Another closely related work is that of Penczek [17, 18] which defines a temporal logic of causal knowledge. Knowledge operators are provided to reason about the local history of a process, as well as about the knowledge it acquires from other processes. However, in order to keep the complexity of model checking tractable, Penczek does not allow the nesting of causal knowledge operators. Interestingly, the nesting of causal knowledge operators does not add any complexity to our algorithm for monitoring.

Leucker investigates linear temporal logic interpreted over restricted labeled partial orders called Mazurkiewicz traces [12]. An overview of distributed linear time temporal logics based on Mazurkiewicz traces is given by Thiagarajan *et al.* in [22]. Alur *et al.* [4] introduce a temporal logic of causality (TLC) which is interpreted over causal structures corresponding to partial order executions of a distributed system. They use both past and future time operators and give a model checking algorithm for the logic.

In recent years, there has been considerable interest in runtime verification [1]. Havelund *et al.* [10] gives algorithms for synthesizing efficient monitors for safety properties. Sen *et al.* [20] develop techniques for runtime safety analysis for multithreaded programs and introduce the tool JMPAX. Some other runtime verification systems include JPaX from NASA Ames [9] and UPENN’s Mac [11].

We can think of at least three major contributions of the work presented in this paper. First, we define a simple but expressive logic to specify safety properties in distributed systems. Second, we provide an algorithm to synthesize decentralized monitors for safety properties that are expressed in the logic. Finally, we describe the implementation of a tool (DIANA) that is based on this technique. The tool is publicly available for download.

The rest of the paper is organized as follows. Section 2 and Section 3 gives the preliminaries. Section 4 introduces PT-DTL. In Section 5 we describe the algorithm that underlies our implementation. Section 6 briefly describes the implementation along with initial experimentation.

2. Distributed Systems

We consider a distributed system as a collection of n processes (p_1, \dots, p_n) , each having a local state and communicating with each other through asynchronous message exchange. The computation of each process is abstracted out in terms of *events*, while the distributed computation is abstracted out in terms of a partial order \prec on events. There can be three types of events:

1. *internal*: an event that only changes the local state of a process by changing the values of its local variables;
2. *send*: an event denoting the sending of a message by a process to another process, and
3. *receive*: an event denoting the reception of a message by a process from another process.

Let E_i denote the set of events of process p_i and let E denote $\bigcup_i E_i$. Also, let $\lessdot \subseteq E \times E$ be defined as follows.

1. $e \lessdot e'$ if e and e' are events of the same process and e happens immediately before e' ,
2. $e \lessdot e'$ if e is the send event of a message at some process and e' is the corresponding receive event of the message at the recipient process.

The partial order \prec is the transitive closure of the relation \lessdot . This partial order captures the *causality* relation among the events in different processes. The structure described by $\mathcal{C} = (E, \prec)$ is called a *distributed computation*. In what follows, we assume an arbitrary but fixed distributed computation \mathcal{C} . Let us define \preccurlyeq as the reflexive and transitive closure of \lessdot . In Fig. 2, $e_{11} \prec e_{23}$, $e_{12} \prec e_{23}$, and $e_{11} \preccurlyeq e_{23}$. However, $e_{12} \not\preccurlyeq e_{23}$.

The *local state* of a process is abstracted out in terms of a set of events. For $e \in E$ we define $\downarrow e \stackrel{\text{def}}{=} \{e' \mid e' \preccurlyeq e\}$, that is, $\downarrow e$ is the set of events that causally precede e . For $e \in E_i$, we can think of $\downarrow e$ as the local state of p_i when the event e has just occurred. This state contains the history of events of all processes that causally precede e .

We extend the definition of \lessdot , \prec and \preccurlyeq to local states such that $\downarrow e \lessdot \downarrow e'$ iff $e \lessdot e'$, $\downarrow e \prec \downarrow e'$ iff $e \prec e'$, and $\downarrow e \preccurlyeq$

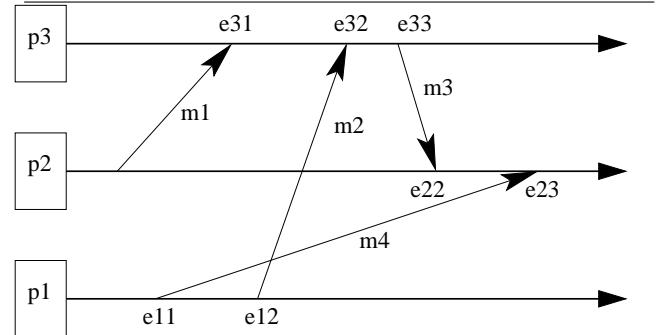


Figure 2. Sample Distributed Computation

$\downarrow e'$ iff $e \preccurlyeq e'$. We denote the set of local states of a process p_i by $LS_i \stackrel{\text{def}}{=} \{\downarrow e \mid e \in E_i\}$ and let $LS \stackrel{\text{def}}{=} \bigcup_i LS_i$. We use the symbols s_i, s'_i, s''_i and so on to represent the local states of process p_i . We also assume that each local state s_i of each process p_i associates values to some local variables V_i , and that $s_i(v)$ denotes the value of a variable $v \in V_i$ in the local state s_i at process p_i .

We use the notation $causal_j(s_i)$ to refer to the latest state of process p_j of which process p_i knows while in state s_i . Formally, if $causal_j(s_i) = s_j$ then $s_j \in LS_j$ and $s_j \preccurlyeq s_i$ and for all $s'_j \in LS_j$ if $s'_j \preccurlyeq s_i$ then $s'_j \preccurlyeq s_j$. For example, in Figure 2 $causal_1(\downarrow e_{23}) = \downarrow e_{12}$. Note that if $i = j$ then $causal_j(s_i) = s_i$.

3. Past Time Linear Temporal Logic (PT-LTL)

Past-time Linear Temporal Logic (PT-LTL) [13, 14] has been successfully used in [10, 11, 20] to express, monitor and predict violations of safety properties in software systems. The syntax of PT-LTL is given as follows:

$$\begin{aligned} F ::= & \text{ true } | \text{ false } | a \in A | \neg F | F \text{ op } F & \text{ propositional} \\ & | \odot F | \diamond F | \square F | F \mathcal{S} F & \text{ temporal} \end{aligned}$$

where op are the standard binary operators, namely $\wedge, \vee, \rightarrow, \leftrightarrow$, and $\odot F$ should be read as “previously”, $\diamond F$ as “eventually in the past”, $\square F$ as “always in the past”, $F_1 \mathcal{S} F_2$ as “ F_1 since F_2 ”.

The logic is interpreted on a finite sequence of states or a run. If $\rho = s_1 s_2 \dots s_n$ is a run then we let ρ_i denote the prefix run $s_1 s_2 \dots s_i$ for each $1 \leq i \leq n$. The semantics of the different operators is given in Table 1.

For example, the formula $\square((action \wedge \odot \neg action) \rightarrow (\neg stop \mathcal{S} start))$ states that whenever *action* starts to be true, it is the case that *start* was true at some moment in the past and since then *stop* was never true, or in other words that the action is taken only when the system is active.

Notice that the semantics of “previously” is given as if the trace is unbounded in the past and stationary in the first event. In runtime monitoring, we start the process of moni-

$\rho \models \text{true}$	for all ρ ,
$\rho \not\models \text{false}$	for all ρ ,
$\rho \models a$	iff a holds in the state s_n ,
$\rho \models \neg F$	iff $\rho \not\models F$,
$\rho \models F_1 \text{ op } F_2$	iff $\rho \models F_1$ and/or/implies/iff $\rho \models F_2$, when op is $\wedge / \vee / \rightarrow / \leftrightarrow$,
$\rho \models \odot F$	iff $\rho' \models F$, where $\rho' = \rho_{n-1}$ if $n > 1$ and $\rho' = \rho$ if $n = 1$,
$\rho \models \Diamond F$	iff $\rho_i \models F$ for some $1 \leq i \leq n$,
$\rho \models \Box F$	iff $\rho_i \models F$ for all $1 \leq i \leq n$,
$\rho \models F_1 \mathcal{S} F_2$	iff $\rho_j \models F_2$ for some $1 \leq j \leq n$ and $\rho_i \models F_1$ for all $j < i \leq n$,

Table 1. Semantics of PT-LTL

toring from the point the first event is generated and it continues as long as the events are generated.

The logic PT-LTL is interpreted over a linear execution trace. However, in distributed systems a computation is represented by a partial order which can have several possible linearizations. Therefore, monitoring a distributed computation requires monitoring of all the possible linear traces that can be obtained from the partial order. The number of linearizations of a partial order can be exponential in the length of the computation. This makes monitoring of PT-LTL over a distributed computation intractable. A major contribution of this paper is to extend PT-LTL so that we can reason about a distributed property using only local monitoring. We describe this extension in the next section.

4. Past Time Distributed Temporal Logic

In this section we extend PT-LTL to a new logic, called past-time Distributed Temporal Logic (PT-DTL), to express safety properties of distributed message passing systems. Although PT-LTL works well for a single process, once we have more processes interacting with each other we need to reason about the state of remote processes. Since practical distributed systems are usually asynchronous and the absolute global state of the system is *not* available to processes, the best thing that each process can do it to reason about the global state that it is *aware of*.

In order to reason about the global distributed computation locally, we add a pair of new operators, called *epistemic operators* as in [19], written @ , whose role is to evaluate an expression or a formula in the *last known state* of a remote process. We call such an expression or a formula *remote*. A remote expression or formula can contain nested epistemic operators and refers to variables that are local to the remote process. By allowing remote expressions in addition to remote formulae, we allow one to specify a larger class of desirable properties of distributed systems without sacrificing the efficiency of monitoring.

Consider, for example, the simple local property at a process p_i that if α is true in the current local state of p_i then β must be true at the latest state of process p_j of which

p_i is aware of. This will be written formally in PT-DTL as $\alpha \rightarrow @_j \beta$. However, referring to remote formulae only is *not* sufficient in order to express a broad range of useful global properties such as “at process p_i , the value of x in the current state is greater than the value of y at process p_j in the latest causally preceding state”. The reason we introduce the novel epistemic operators on expressions is that, in order to state many properties of interest in distributed systems, we find it crucial to be able to also refer to *values* of expressions in remote local states. For example, the property above can be formally specified as the PT-DTL formula $x > @_j y$ at process p_i . Here $@_j y$ is the value of y at process p_j that p_i is aware of.

The intuition underlying PT-DTL is that each process is associated with local temporal formulae which, due to the epistemic operators, can refer to the global state of the distributed system. These formulae are required to be valid at the respective processes during a distributed computation. The distributed computation satisfies the specification when all the local formulae are shown to satisfy the computation. Next we formally describe the syntax and semantics of PT-DTL.

4.1. Syntax

In the sequel, whenever we talk about a PT-DTL formula, it is in the context of a particular process, say p_i . We call such formulae *i-formulae* and let them be denoted by F_i, F'_i , etc. Additionally, we introduce the notion of expressions local to a process p_i called as *i-expressions* and let them be denoted by ξ_i, ξ'_i , etc. Informally, an *i-expression* is an expression over the global state of the system that process p_i is currently aware of. Local predicates on *i-expressions* form the atomic propositions on which the temporal *i-formulae* are built.

We add the *epistemic operators* $@_j$ that take j -expressions or j -formulae and convert them into expressions or formulae local to process p_i . Informally, $@_j$ yields an expression or a formula on process p_j over the projection of the global state that the current process is aware of. The following gives the formal syntax of PT-DTL with re-

spect to a process p_i , where i and j are any process indexes (not necessarily distinct):

$$\begin{aligned}
 F_i ::= & \text{ true } | \text{ false } | P(\vec{\xi}_i) | \neg F_i | F_i \text{ op } F_i & \text{propositional} \\
 & | \odot F_i | \diamond F_i | \Box F_i | F_i \mathcal{S} F_i & \text{temporal} \\
 \xi_i ::= & c | v_i | f(\vec{\xi}_i) & \text{functional} \\
 & | @_j F_j & \text{epistemic} \\
 \vec{\xi}_i ::= & (\xi_i, \dots, \xi_i) & \text{epistemic}
 \end{aligned}$$

The infix operator op can be any binary propositional operator such as $\wedge, \vee, \rightarrow, \equiv$. The term $\vec{\xi}_i$ stands for a tuple of expressions on process p_i . The term $P(\vec{\xi}_i)$ is a (computable) predicate over the tuple $\vec{\xi}_i$ and $f(\vec{\xi}_i)$ is a (computable) function over the tuple. For example, P can be $<, \leq, >, \geq, =$. Similarly, some examples of f are $+, -, /, *$. Variables v_i belongs to the set V_i which contains all the local state variables of process p_i . Constants such as $0, 1, 3.4$ are represented by c .

The expression $@_j \xi_j$ is an i -expression representing the remote expression ξ_j . Similarly, $@_j F_j$ is an i -formula referring to the local knowledge about the remote validity of j -formula F_j . In other words, $@_j$ converts a j -expression or a j -formula to an i -expression or an i -formula, respectively.

4.2. Semantics

The semantics of PT-DTL is a natural extension of PT-LTL with epistemic operators. The atomic propositions of PT-LTL are replaced by predicates over tuples of expressions. Table 2 formally gives the semantics of each operator of PT-DTL. $(C, s_i)[@_j \xi_j]$ is the value of the expression ξ_j in the state $s_j = \text{causal}(s_i)$ which is the latest state of process p_j of which process p_i is aware of. We assume that expressions are properly typed. Typically these types would be: `integer`, `real`, `strings` etc. We also assume that $s_i, s'_i, s''_i, \dots \in LS_i$ and $s_j, s'_j, s''_j, \dots \in LS_j$. Notice that, like in PT-LTL, the meaning of the “previously” operator on the initial state of each process reflects the intuition that the execution trace is unbounded in the past and stationary. We consider this as the most reasonable assumption that one can make about the past.

4.3. Examples

To demonstrate the power and expressiveness of PT-DTL, we consider several relatively standard examples in the distributed systems literature (see, e.g., [21]).

The first example is regarding the well known problem of leader election for a network of processes. The key requirement for leader election is that there is at-most one leader. If there are n processes and `state` is a variable in each process that can have values `leader`, `loser`, `candidate`, `sleep`, then we can write the property at every process as: “if a leader is elected

then if the current process is a leader then, to its knowledge, none of the other processes is a leader”. We can formalize this requirement as the following PT-DTL i -local formula:

$$\text{leaderElected} \rightarrow (\text{state} = \text{leader} \rightarrow \bigwedge_{j \neq i} (@_j(\text{state} \neq \text{leader}))$$

Given an implementation of the leader election problem, one can monitor this formula at every process p_i . If violated then clearly the leader election implementation is incorrect.

The second example is concerned with a number of processes which vote on a particular resolution. The desired property, “if the resolution is accepted then more than half of the processes say yes”, can be stated as:

$$\text{accepted} \rightarrow (@_1(\text{vote}) + @_2(\text{vote}) + \dots + @_n(\text{vote})) > n/2$$

Here, a process stores 1 in a local variable `vote` if it is in favor of the resolution, and 0 otherwise.

The third example is a safety property that a server must satisfy in case it reboots itself. The property is that “server accepts to reboot only after knowing that each client is inactive and aware of the warning to reboot”. The property can be written as a *server*-local formula as follows:

$$\text{rebootAccepted} \rightarrow \bigwedge_{\text{client}} (@_{\text{client}}(\text{inactive} \wedge @_s \text{rebootWarning}))$$

Note that the above formula contains nested epistemic operators.

5. Monitoring Algorithm for PT-DTL

In this section, we describe an automated technique to synthesize efficient distributed monitors for safety properties in distributed systems expressed in PT-DTL. We assume that one or more processes are associated PT-DTL formulae which must be satisfied by the distributed computation. The synthesized monitor is *distributed*, in the sense that it consists of separate, *local monitors* running on each process. A local monitor can attach additional information to any outgoing message from the corresponding process. This information can subsequently be extracted by the monitor on the receiving side without changing the underlying semantics of the distributed program. The key guiding principles in the design of this technique are:

1. The local monitors should be fast, so that monitoring can be done online;
2. The local monitors should have little memory overhead, in particular, it should *not* need to store the entire history of events on a process since this can be quite large;

$\mathcal{C}, s_i \models \text{true}$	for all s_i
$\mathcal{C}, s_i \not\models \text{false}$	for all s_i
$\mathcal{C}, s_i \models P(\xi_i, \dots, \xi'_i)$	iff $P((\mathcal{C}, s_i)[\xi_i], \dots, (\mathcal{C}, s_i)[\xi'_i]) = \text{true}$
$\mathcal{C}, s_i \models \neg F_i$	iff $\mathcal{C}, s_i \not\models F_i$
$\mathcal{C}, s_i \models F_i \text{ op } F'_i$	iff $\mathcal{C}, s_i \models F_i \text{ op } \mathcal{C}, s_i \models F'_i$
$\mathcal{C}, s_i \models \odot F_i$	iff if $\exists s'_i . s'_i \preceq s_i$ then $\mathcal{C}, s'_i \models F_i$ else $\mathcal{C}, s_i \models F_i$
$\mathcal{C}, s_i \models \diamond F_i$	iff $\exists s'_i . s'_i \preceq s_i$ and $\mathcal{C}, s'_i \models F_i$
$\mathcal{C}, s_i \models \Box F_i$	iff $\mathcal{C}, s_i \models F_i$ for all $s'_i \preceq s_i$
$\mathcal{C}, s_i \models F_i \mathcal{S} F'_i$	if $\exists s'_i . s'_i \preceq s_i$ and $\mathcal{C}, s'_i \models F'_i$ and $\forall s''_i . s'_i \prec s''_i \preceq s_i$ implies $\mathcal{C}, s''_i \models F_i$
$\mathcal{C}, s_i \models @_j F_j$	iff $\mathcal{C}, s_j \models F_j$ where $s_j = \text{causal}_j(s_i)$
$(\mathcal{C}, s_i)[v_i]$	$= s_i(v_i)$, that is, the value of v_i in s_i
$(\mathcal{C}, s_i)[c_i]$	$= c_i$
$(\mathcal{C}, s_i)[f(\xi_i, \dots, \xi'_i)]$	$= f((\mathcal{C}, s_i)[\xi_i], \dots, (\mathcal{C}, s_i)[\xi'_i])$
$(\mathcal{C}, s_i)[@_j \xi_j]$	$= (\mathcal{C}, s_j)[\xi_j]$ where $s_j = \text{causal}_j(s_i)$

Table 2. Semantics of PT-DTL

3. Additional messages needed to be sent between processes solely for the purpose of monitoring should be minimized and ideally, should be zero. Further, additional information piggybacked on regular messages (those generated as part of the distributed computation) should be small.

In this section, when we refer to a remote expression or formulae we mean one which occurs in any of the monitored PT-DTL formulae.

5.1. Knowledge Vectors

Consider evaluating a remote j -expression $@_j \xi_j$ at process p_i . The naive solution is that process p_j simply piggybacks the value of ξ_j evaluated at p_j , with every message that it sends out. The recipient process p_i can extract this value and use it as the value of $@_j \xi_j$. However, this simplistic approach is problematic, because messages from p_j could reach p_i in arbitrary order. A message, sent earlier but received later, could supersede the most recent value of $@_j \xi_j$ received in a message that was sent later but received earlier. This can also happen for remote formulae. To keep track of the causal history, or in other words the most recent knowledge, we also need the event number at p_j at which these expressions were sent out in messages so that stale information in a reordered message sequence can be discarded.

Causal ordering can be effectively accomplished by using an array called KNOWLEDGEVECTOR with an entry for any process p_j for which there is an occurrence of $@_j$ in any PT-DTL formula at any process. Knowledge vectors are motivated and inspired by vector clocks [8, 15]. The size of KNOWLEDGEVECTOR is not dependent on the number of processes but on the number of remote expressions and for-

mulae. Let $KV[j]$ denote the entry for process p_j on a vector KV . $KV[j]$ contains the following fields:

- The sequence number of the last event seen at p_j , denoted by $KV[j].seq$;
- A set of values $KV[j].values$ storing the values j -expressions and j -formulae.

Each process p_i keeps a local KNOWLEDGEVECTOR denoted by KV_i . The monitor of process p_i attaches a copy of KV_i with every outgoing message m . We denote the copy by KV_m . The algorithm for the update of KNOWLEDGEVECTOR KV_i at process p_i is given below:

1. **[internal]:** update $KV_i[i]$. For this we evaluate $eval(\xi_i, s_i)$ and $eval(F_i, s_i)$ (described in Subsection 5.2) for each $@_i \xi_i$ and $@_i F_i$, respectively, and store them in the set $KV_i[i].values$;
2. **[send m]:** $KV_i[i].seq \leftarrow KV_i[i].seq + 1$. Send KV_i with m as KV_m ;
3. **[receive m]:** for all j , if $KV_m[j].seq > KV_i[j].seq$ then $KV_i[j] \leftarrow KV_m[j]$, that is, $KV_i[j].seq \leftarrow KV_m[j].seq$, and $KV_i[j].values \leftarrow KV_m[j].values$.

We call this the KNOWLEDGEVECTOR algorithm. Informally, $KV_i[j].values$ contains the latest values that p_i has for j -expressions or j -formulae. Therefore, for the value of a remote expression or formula of the form $@_j \xi_j$ or $@_j F_j$, process p_i can just use the entry corresponding to ξ_j or F_j in the set $KV_i[j].values$. Note that the sequence number needs to be incremented only when sending messages. The correctness of this algorithm is relatively straightforward and therefore we skip its formal proof:

Proposition 1 *For any process p_i and any j , the entry for ξ_j or F_j in $KV_i[j].values$ contains the value of $@_j \xi_j$ or $@_j F_j$, respectively.*

The algorithm above tries to minimize the local work when sending a message. However, notice that the values calculated at step 1 are needed only when an outgoing message is generated at step 2, so one could have just evaluated all the expressions ξ_i and F_i at step 2, right before the message is sent out. This would reduce the runtime overhead at step 1 but it would increase it at step 2. Depending on the specific application under consideration, one may prefer one way or the other.

The initial values for all the variables in the distributed program can be found either by a static analysis of the program or by a distributed broadcast at the beginning of the computation. Thus, it is assumed that each process p_i has complete knowledge of the initial values of remote expressions for all processes. These values are then used to initialize the entries $KV_i[j].values$ in the KNOWLEDGEVECTOR of p_i for all j .

5.2. Monitoring a Local PT-DTL Formula

We now describe the details of a local monitor. The monitoring algorithm for a PT-DTL formula is similar in spirit to that for an ordinary PT-LTL formula described in [20]. The key difference is that we allow remote expressions and remote formulae whose values and validity, respectively, need to be transferred from the remote process to the current process. Once the KNOWLEDGEVECTOR is properly updated, the local monitor can compute the boolean value of the formula to be monitored, by recursively evaluating the boolean value of each of its subformulae in the current state. In order to do that, it may also use the boolean values of subformulae evaluated in the previous state and the values of remote expressions and remote formulae.

A function $eval$ is defined next, which takes advantage of the recursive nature of the temporal operators, as described in Table 3, to calculate the boolean value of a formula in the current state in terms of (a) its boolean value in the previous state and (b) the boolean value of its subformulae in the current state. The function $op(F_i)$ returns the operator of the formula F_i , $binary(op(F_i))$ returns *true* if $op(F_i)$ is binary, $unary(op(F_i))$ returns *true* if $op(F_i)$ is unary, $left(F_i)$ returns the left subformula of F_i , $right(F_i)$ returns the right subformula of F_i when $op(F_i)$ is binary, and $subformula(F_i)$ returns the subformula of F_i otherwise. The variable $index$ represents the index of a subformula :

```
array now; array pre; int index;

boolean eval(Formula F_i, State s_i){
  if binary(op(F_i)) then{
    lval ← eval(left(F_i), s_i);
    rval ← eval(right(F_i), s_i);}
  else if unary(op(F_i)) then
    val ← eval(subformula(F_i), s_i);
  index ← 0;
```

```
case(op(F_i)) of{
  true : return true; false : return false;
  P(ξ̄_i) : return P(eval(ξ_i, s_i), ..., eval(ξ'_i, s_i)));
  op : return rval op lval; ¬ : return not val;
  S : now[index] ← (pre[index] and lval) or rval;
  return now[index++];
  □ : now[index] ← pre[index] and val;
  return now[index++];
  ◊ : now[index] ← pre[index] or val;
  return now[index++];
  ⊖ : now[index] ← val; return pre[index++];
  @_j F_j : return value of F_j from KV_i[j].values;
}
```

Here, the global array pre contains the boolean values of all subformulae in the previous state that will be required in the current state, while the global array now , after the evaluation of $eval$, will contain the boolean values of all subformulae in the current state that may be required in the next state. Note that the now array's value is set in the function $eval$. The function $eval$ on expressions is defined next:

```
value eval(Expression ξ_i, State s_i){
  case(ξ_i) of{
    v_i: return s_i(v_i); c_i: return c_i;
    f(ξ_i^1, ..., ξ_i^k): return f(eval(ξ_i^1, s_i), ..., eval(ξ_i^k, s_i));
    @_j ξ'_j: return value of ξ'_j from KV_i[j].values;
  }
}
```

The function $eval$ cannot be used to evaluate the boolean value of a formula at the first event, as the recursion handles the case $n = 1$ in a different way. We define the function $init$ to handle this special case as implied by the semantics of PT-DTL in Tables 2 and 3 on one event traces:

```
boolean init(Formula F_i, State s_i){
  if binary(op(F_i)) then{
    lval ← init(left(F_i, s_i));
    rval ← init(right(F_i, s_i));}
  else if unary(op(F_i)) then
    val ← init(subformula(F_i, s_i));
  index ← 0;
  case(op(F_i)) of{
    true : return true; false : return false;
    P(ξ̄_i) : return P(eval(ξ_i, s_i), ..., eval(ξ'_i, s_i)));
    op : return rval op lval; ¬ : return not val;
    S : now[index] ← rval; return now[index++];
    □, ◊, ⊖ : now[index] ← val; return now[index++];
  }
}
```

As mentioned earlier in the KNOWLEDGEVECTOR algorithm, we either use the function $eval$ after every internal event (updates of local variables that are referred in the formula) or immediately before sending any message, in or-

$$\begin{aligned}
\mathcal{C}, s_i \models \Diamond F_i &= \mathcal{C}, s_i \models F_i \text{ or } (\exists s'_i \cdot s'_i \lessdot s_i \text{ and } \mathcal{C}, s'_i \models \Diamond F_i) \\
\mathcal{C}, s_i \models \Box F_i &= \mathcal{C}, s_i \models F_i \text{ and } (\exists s'_i \cdot s'_i \lessdot s_i \text{ implies } \mathcal{C}, s'_i \models \Box F_i) \\
\mathcal{C}, s_i \models F_i \mathcal{S} F'_i &= \mathcal{C}, s_i \models F'_i \text{ or } \\
&(\mathcal{C}, s_i \models F_i \text{ and } \exists s'_i \cdot s'_i \lessdot s_i \text{ and } \mathcal{C}, s'_i \models F_i \mathcal{S} F'_i)
\end{aligned}$$

Table 3. Recursive Semantics of PT-DTL

der to properly update the set $KV[i].values$. We then assign *now* to *pre*. If a monitored PT-DTL formula F_i is specified for a process p_i , we call p_i as the owner of that formula. At the owner, process we evaluate F_i using the *eval* function after every internal and receive event and assign *now* to *pre*. This is done after the KNOWLEDGEVECTOR is updated correspondingly after the event. If the evaluation of F_i is false then we report a warning that the formula F_i is violated.

The time and space complexity of this algorithm at every event is $\Theta(m)$, where m is the size of the original local formula.

5.3. Example

Let us consider three processes, p_1 , p_2 and p_3 . Process p_1 has a local variable x whose initial value is 5. Process p_2 has a local variable y with initial value 7. Process p_2 monitors the formula $\Box(y \geq @_1 x)$. An example computation is shown in Figure 3.

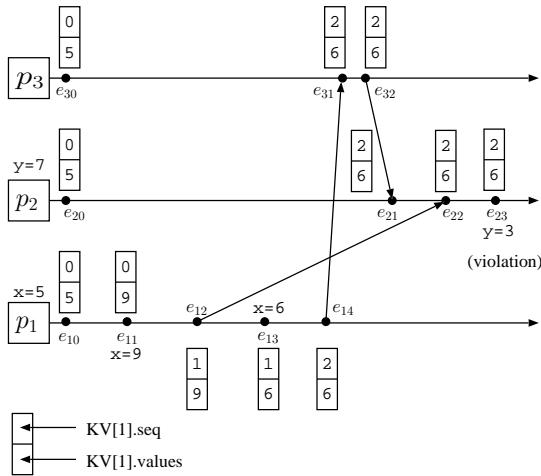


Figure 3. Monitoring of $\Box(y \geq @_1 x)$ at p_2

There is only one formula to monitor with a single occurrence of an $@$ operator, namely $@_1 x$. Hence, the KNOWLEDGEVECTOR has a single entry which corresponds to p_1 . Moreover, since the only remote expression to be tracked is x , $KV[1].values$ simply stores the value of x . In the figure, next to each event, we show $KV[1]$ at that instant for

that process. $KV[1]$ is graphically displayed by a stack of two numbers, the top number showing $KV[1].seq$ and the bottom number showing the value for x .

The computation starts off with the initial values of $x = 5$ and $y = 7$. All processes know the initial value of x , hence the $KV[1].values$ for each process has value 5. It is easy to see that the monitored formula $\Box(y \geq @_1 x)$ holds initially at p_2 . Subsequently, at p_1 there is an internal event e_{11} which sets $x = 9$ and updates $KV_1[1].values$ correspondingly. Process p_1 then sends a message to p_2 with a copy of its current KV . Another internal event e_{13} causes x to be set to 6. Process p_1 again sends a message, this time to p_3 , with the updated KV . Process p_3 updates its KV and sends this on the message it sends to p_2 .

At process p_2 , the message sent by p_3 happens to arrive earlier than the message from p_1 . Therefore, at event e_{21} , on receiving the message from p_3 , process p_2 is able to update its KV to the one sent at event e_{14} . The monitor at p_2 again evaluates the property and finds that it still holds. The message sent by p_1 finally arrives at e_{22} but the KV piggybacked on is ignored as it has a smaller $KV[1].seq$ than $KV_2[1].seq$. The monitor correctly continues to declare the property valid. However, another internal event at p_2 causes the value of y to drop to 3, at which point the monitor detects a property violation.

6. The DIANA Tool

We have implemented the above technique as a tool, called DIANA (Distributed ANALysis), whose architecture is depicted in Figure 1. DIANA is publicly available and can be downloaded from its web-site: <http://fsl.cs.uiuc.edu/diana/>. Both DIANA and the framework under which it operates are written in Java.

6.1. The Message Passing Framework

Distributed systems with asynchronous processes can be reasoned about and implemented in a number of different formalisms. We have chosen to use Actors [2, 3] for our system. Actors are a model of distributed reactive objects and have a built-in notion of encapsulation and interaction, making them well suited to represent evolution and coordination among interacting components in distributed applications. Conceptually, an actor encapsulates a state, a thread of control, and a set of procedures which manipulate the

state. Actors coordinate by asynchronously sending messages to each another.

In the actor framework, a distributed system consists of different actors communicating through messages. Thus, there is an actor for each process in the system.

In the implementation, each type of actor (or process) is denoted by a Java class that extends a base class `Actor`. This base class implements a message queue and provides the method `send` for asynchronous message sending. Each actor object executes in a separate thread. The state of an actor is represented by the fields of the Java class. Each Java class also contains a set of `public` methods that can be invoked in response to messages received from other actors.

The underlying system, which we call *ActorManager*, takes a message and transfers it to the message queue of the target actor. The target actor takes an available message from the message queue and invokes the method mentioned in the message. While processing a message, an actor may send messages to other actors. Message sending, being asynchronous, never blocks an actor when it sends a message. However, it blocks if there is no message in its message queue. The system is initialized by the *ActorManager* object that creates all the actors in the system and starts the execution of the system.

6.2. Distributed Monitors in DIANA

The user of DIANA specifies the local PT-DTL formulae to be monitored on each actor in a special file. Each actor has a unique name, which is the name of the corresponding process. The name is passed as a string at the time creation of an actor.

As Figure 1 shows, the core of DIANA consists of two modules: an *instrumentation* module and a *monitoring* module. The instrumentation module takes the specification file and the distributed program written in the above framework and creates a Java class `MonitorImpl` that implements a local monitor for each actor (or process). It also automatically instruments the distributed program *at the bytecode level* (after compilation), so that the distributed program invokes its local monitor whenever it modifies a field variable (internal event), sends a message, or invokes a method (receive event).

One can also choose to evaluate the epistemic expressions and formulae immediately before sending an event as explained in Subsection 5.1, in which case the local monitor is not invoked when field variables are modified. While runtime overhead is not a major concern for us at this stage, this tends to reduce the runtime overhead in most situations.

6.3. Test Cases

We have implemented a voting algorithm in the framework above. In this algorithm, a `Chair` process asks

for vote on a resolution from N voters named `Voter1`, `Voter2`, ..., `VoterN`, where N is initialized to an arbitrary but fixed positive number. We assumed that the processes were connected in a tree kind of network with the `Chair` at the root of the tree and the voters at different nodes. Each voter randomly decides if it wants to vote in favor or against the resolution and stores 1 or 0 respectively in a local state variable `vote`. It then sends the decision to its immediate parent in the tree. The parent collects the votes and sends the sum of its vote and its progenies' votes to its immediate parent. The `Chair` process collects all the votes and rejects the resolution only if half or more voters have rejected. We monitor the following safety property at `Chair`:

$$\text{reject} \rightarrow ((\sum_{i \in [1..N]} @_{\text{Voter}_i}(\text{vote})) < N/2)$$

The property was found to be violated at several runs. The reason was that at some voter nodes, the voter sent the sum of its progenies' votes without adding its own vote. This resulted in the rejection of the resolution when it should have been accepted.

We have also tested a vector clock [8, 15] algorithm implemented in the framework presented in this section. The algorithm was implemented as part of global snapshot and garbage collection algorithm. In this algorithm, every process is assumed to have a local vector clock V that it updates according to the standard vector clock algorithm [8] whenever there is an internal event, a send event or a receive event. The safety property that this must satisfy is that at every process p_i , “all entries of the local vector clock must be greater than or equal to the local vector clock in a causally latest preceding state of any other process”. This property can be expressed as the following *i*-formula:

$$\square(\bigwedge_{j \in [1..n]} V \geq @_j V)$$

where $V \geq V'$ when every entry in V is greater than or equal to the corresponding entry in V' . Another safety property states that “at every process p_i the i -th entry in its local vector clock must be strictly greater than the i -th entry of the local vector clock of any other process”. This can be expressed as the following *i*-formula:

$$\square(\bigwedge_{j \in [1..n]} V[i] > @_j V[i])$$

The second property was found to be violated in some computations due to a bug caused by failure to increment the i -th entry of the local vector clock of process p_i when receiving events.

These simple examples show the practical use and power of PT-DTL and the monitoring tool DIANA based on it.

7. Conclusion and Future Work

This work represents the first step in effective distributed monitoring. The work presented here suggests a number of problems that require further research. The logic itself could be made more expressive so that it expresses not only safety, but also liveness properties. One difficulty is that software developers are reluctant to use formal notations. A partial solution may be to merge the present work with a more expressive and programmer friendly monitoring temporal logic such as EAGLE [6]. A complementary approach is to develop visual notations and synthesizing temporal logic formulas from such notations. There may also be the possibility of learning formulas based on representative scenarios.

An interesting avenue of future investigation that our work suggests is what we call *Knowledge-based Aspect-Oriented Programming*. Knowledge-based Aspect-Oriented Programming is a meta-programming discipline that is suitable for distributed applications. In this programming paradigm, appropriate actions are associated with each safety formula; these actions are taken whenever the formula is violated to guide the program and avoid catastrophic failures.

Acknowledgements

The first three authors are supported in part by the Defense Advanced Research Projects Agency (the DARPA IPTO TASK Program, contract F30602-00-2-0586, the DARPA IXO NEST Program, contract F33615-01-C-1907), the ONR Grant N00014-02-1-0715, and the Motorola Grant MOTOROLA RPS #23 ANT. The last author is supported in part by the joint NSF/NASA grant CCR-0234524.

References

- [1] *1st, 2nd and 3rd CAV Workshops on Runtime Verification (RV'01 - RV'03)*, volume 55(2), 70(4), 89(2) of *Electronic Notes in Theoretical Computer Science*. Elsevier Science: 2001, 2002, 2003.
- [2] G. Agha. *Actors: A Model of Concurrent Computation*. MIT Press, 1986.
- [3] G. Agha, I. A. Mason, S. F. Smith, and C. L. Talcott. A foundation for actor computation. *Journal of Functional Programming*, 7:1–72, 1997.
- [4] R. Alur, D. Peled, and W. Penczek. Model checking of causality properties. In *Proceedings of the 10th Annual IEEE Symposium on Logic in Computer Science (LICS'95)*, pages 90–100, San Diego, California, 1995.
- [5] R. Aumann. Agreeing to disagree. *Annals of Statistics*, 4(6):1236–1239, 1976.
- [6] H. Barringer, A. Goldberg, K. Havelund, and K. Sen. Rule-based runtime verification. In *Proceedings of 5th International Conference on Verification, Model Checking and Abstract Interpretation (VMCAI'04)*, volume 2937 of *Lecture Notes in Computer Science*, pages 44–57, Venice, Italy, January 2004. Springer-Verlag.
- [7] R. Fagin, J. Halpern, Y. Moses, and M. Vardi. *Reasoning about Knowledge*. MIT Press, 1995.
- [8] C. J. Fidge. Partial orders for parallel debugging. In *Proceedings of the 1988 ACM SIGPLAN and SIGOPS Workshop on Parallel and Distributed Debugging (WPDD'88)*, pages 183–194. ACM, 1988.
- [9] K. Havelund and G. Roşu. Java pathexplorer – A runtime verification tool. In *The 6th International Symposium on Artificial Intelligence, Robotics and Automation in Space: A New Space Odyssey*, Montreal, Canada, June 18 - 21, 2001.
- [10] K. Havelund and G. Roşu. Synthesizing monitors for safety properties. In *Tools and Algorithms for Construction and Analysis of Systems (TACAS'02)*, volume 2280 of *Lecture Notes in Computer Science*, pages 342–356. Springer-Verlag, 2002.
- [11] M. Kim, S. Kannan, I. Lee, and O. Sokolsky. Java-mac: a run-time assurance tool for java. In *Proceedings of Runtime Verification (RV'01)*, volume 55 of *Electronic Notes in Theoretical Computer Science*. Elsevier Science, 2001.
- [12] M. Leucker. Logics for mazurkiewicz traces. Technical Report AIB-2002-10, RWTH, Aachen, Germany, April 2002.
- [13] Z. Manna and A. Pnueli. *The Temporal Logic of Reactive and Concurrent Systems*. Springer-Verlag, New York, 1992.
- [14] Z. Manna and A. Pnueli. *Temporal Verification of Reactive Systems: Safety*. Springer-Verlag, New York, 1995.
- [15] F. Mattern. Virtual time and global states of distributed systems. In M. C. et. al., editor, *Proceedings of the International Workshop on Parallel and Distributed Algorithms*, pages 215–226. Elsevier Science, 1989.
- [16] B. Meenakshi and R. Ramanujam. Reasoning about message passing in finite state environments. In *International Colloquium on Automata, Languages and Programming (ICALP'00)*, volume 1853 of *Lecture Notes in Computer Science*, pages 487–498. Springer-Verlag, 2000.
- [17] W. Penczek. A temporal approach to causal knowledge. *Logic Journal of the IGPL*, 8(1):87–99, 2000.
- [18] W. Penczek and S. Ambroszkiewicz. Model checking of causal knowledge formulas. In *Workshop on Distributed Systems (WDS'99)*, volume 28 of *Electronic Notes in Theoretical Computer Science*. Elsevier Science, 1999.
- [19] R. Ramanujam. Local knowledge assertions in a changing world. In *Theoretical Aspects of Rationality and Knowledge (TARK'96)*, pages 1–14. Morgan Kaufmann, 1996.
- [20] K. Sen, G. Roşu, and G. Agha. Runtime safety analysis of multithreaded programs. In *ACM SIGSOFT Conference on the Foundations of Software Engineering / European Software Engineering Conference (FSE / ESEC'03)*, Helsinki, Finland, 2003.
- [21] G. Tel. *Introduction to Distributed Algorithms*. Cambridge University Press, 2nd edition, September 2000.
- [22] P. S. Thiagarajan and I. Walukiewicz. An expressively complete linear time temporal logic for Mazurkiewicz traces. In *Twelfth Annual IEEE Symposium on Logic in Computer Science (LICS'97)*, pages 183–194, Warsaw, Poland, 1997.